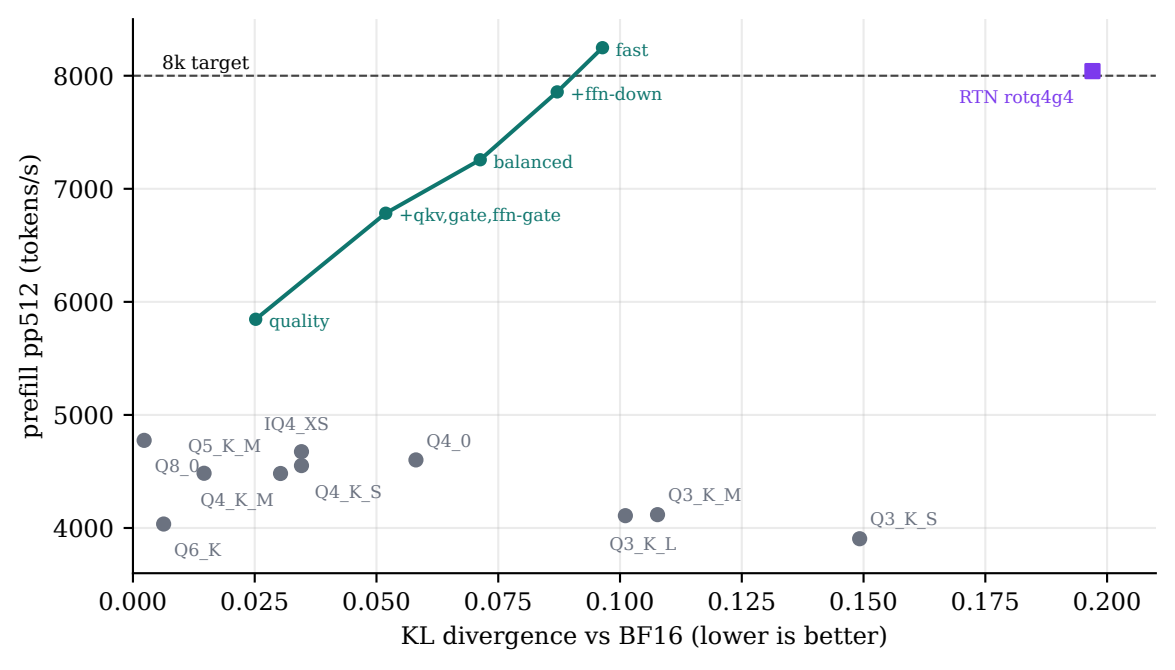
8k target
fast
+ffn-down
balanced
+qkv,gate,ffn-gate
quality
RTN rotq4g4
IQ4_XS
Q5_K_M
Q8_0
Q4_0
Q4_K_M
Q4_K_S
Q3_K_M
Q3_K_S
Q3_K_L
Q6_K
8000
7000
6000
5000
4000
prefill pp512 (tokens/s)
0.000
0.025
0.050
0.075
0.100
0.125
0.150
0.175
0.200
KL divergence vs BF16 (lower is better)
stock llama.cpp quants (ROCm 10.1)
rotated INT4 + GPTQ (this work), per-kind A4 sweep
earlier file: RTN, A4 everywhere