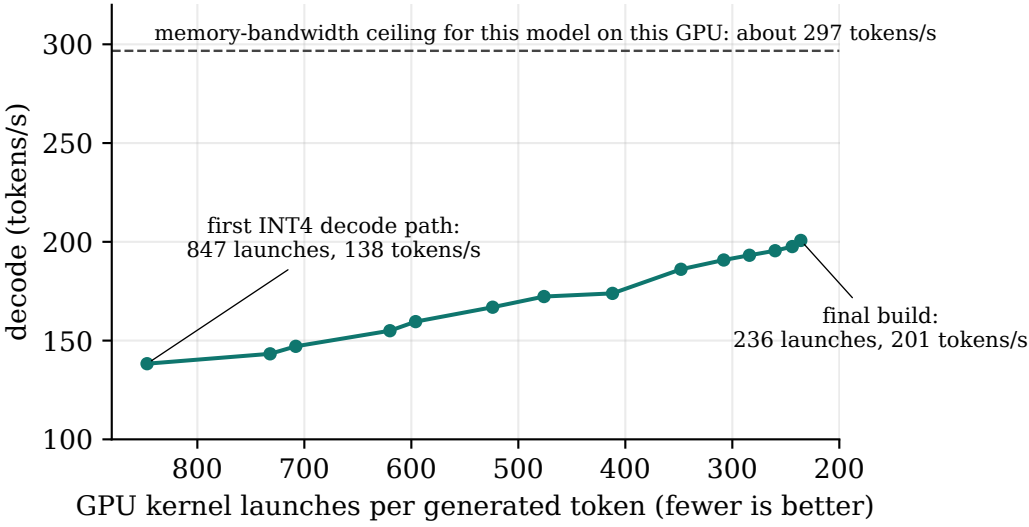
memory-bandwidth ceiling for this model on this GPU: about 297 tokens/s
300
250
200
150
100
decode (tokens/s)
first INT4 decode path:
847 launches, 138 tokens/s
final build:
236 launches, 201 tokens/s
800
700
600
500
400
300
200
GPU kernel launches per generated token (fewer is better)